

IFT 6390

Fondements de l'apprentissage machine

Unsupervised learning

Dimensionality reduction
Continuous latent variables

Professor: Ioannis Mitliagkas
Slides: Pascal Vincent

Context

Supervised Learning = classification, regression

Unsupervised Learning = Algorithms that don't need any explicit "target"/label in the training set.

- Density estimation
(ex: *Gaussian mixture*)
- *Clustering*
(ex: *k-means*)
- Dimension reduction

Dimensionality reduction

What is it?

$$\mathbf{z} \in \mathbb{R}^M$$

$$M < D$$

(0.32, -1.3, 1.2)



$$\mathbf{x} \in \mathbb{R}^D \quad (3.5, -1.7, 2.8, -3.5, -1.4, 2.4, 2.7, 7.5, -3, -2)$$

Dimensionality reduction

What for?

- **Compressing** data (with loss)
- **Visualize** data in 2D or 3D
- **Feature extraction**
potentially +fundamental, +explicative, +compact
Pre-processing => better data representation for another algorithm (classification or regression).

Algorithms

Gaussian linear models

- Principal Component Analysis (PCA)
- Probabilistic PCA
- Factor Analysis

Non-linear/Non-Gaussian models

- Kernel PCA
- Independent Components Analysis
- Auto-encoder neural networks
- Non-linear manifold modeling


Principal Component Analysis (PCA)

- PCA finds a **linear subspace** that is close to the data: **orthogonal projection** of $\mathbf{x} \in \mathbb{R}^D$ onto a linear subspace of lower dimension M .
- Entries of $\mathbf{z} \in \mathbb{R}^M$ are the coordinates of the projection of \mathbf{x} in this M dimensional subspace.
- **Often used** as a pre-processing step (feature extraction) or for visualization.
- An old algorithm with two equivalent formulations:
 - **minimizing the reconstruction error** (Pearson 1901)
 - **maximizing the variance** (Hotelling 1933).
- Probabilistic interpretation is more recent

PCA: two equivalent formulations

We seek the **principal directions** u : an orthonormal basis of the space onto which to project x

- **Minimal reconstruction error:** minimize the mean squared distances between the x 's and their **projections** (blue lines).
- **Maximal variance:** maximize the variance along the projection space (variance of the **green points**)



PCA with variance maximization

Let $\bar{\mathbf{x}} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$ be the empirical mean.

The variance of the points projected onto the direction \mathbf{u}_1 is given by

$$\frac{1}{N} \sum_{n=1}^N (\mathbf{u}_1^T \mathbf{x}_n - \mathbf{u}_1^T \bar{\mathbf{x}})^2 = \mathbf{u}_1^T \mathbf{S} \mathbf{u}_1$$

where \mathbf{S} is the empirical covariance matrix

$$\mathbf{S} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \bar{\mathbf{x}})(\mathbf{x}_n - \bar{\mathbf{x}})^T$$

Adding constraint $\|\mathbf{u}_1\| = 1$, we want to solve the minimization problem

$$\min_{\|\mathbf{u}_1\|=1} \mathbf{u}_1^T \mathbf{S} \mathbf{u}_1$$

whose solution is the dominant eigenvector of the covariance matrix \mathbf{S} .

PCA with variance maximization

- Once we have the first **principal direction/component**, we can incrementally find the second one, third one, etc., by imposing each direction to be orthogonal to the previous ones.
- Using the *minimal reconstruction error* formulation leads to the exact same principal components.

PCA: simple overall procedure

The orthonormal basis \mathbf{U} we obtain contains the M first eigenvectors of the empirical covariance matrix \mathbf{S} (the ones corresponding to the M largest eigenvalues).

Estimate the covariance matrix: $\mathbf{S} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \bar{\mathbf{x}})(\mathbf{x}_n - \bar{\mathbf{x}})^T$

Eigenvalue decomposition:

$$\mathbf{S} = \mathbf{U}\mathbf{\Lambda}\mathbf{U}^T$$

Keep only the first M eigenvectors: $\mathbf{S} \approx \mathbf{U}_M \mathbf{\Lambda}_M \mathbf{U}_M^T$

$$\mathbf{\Lambda} = \begin{pmatrix} \lambda_1 & & 0 \\ & \ddots & \\ 0 & & \lambda_d \end{pmatrix}$$

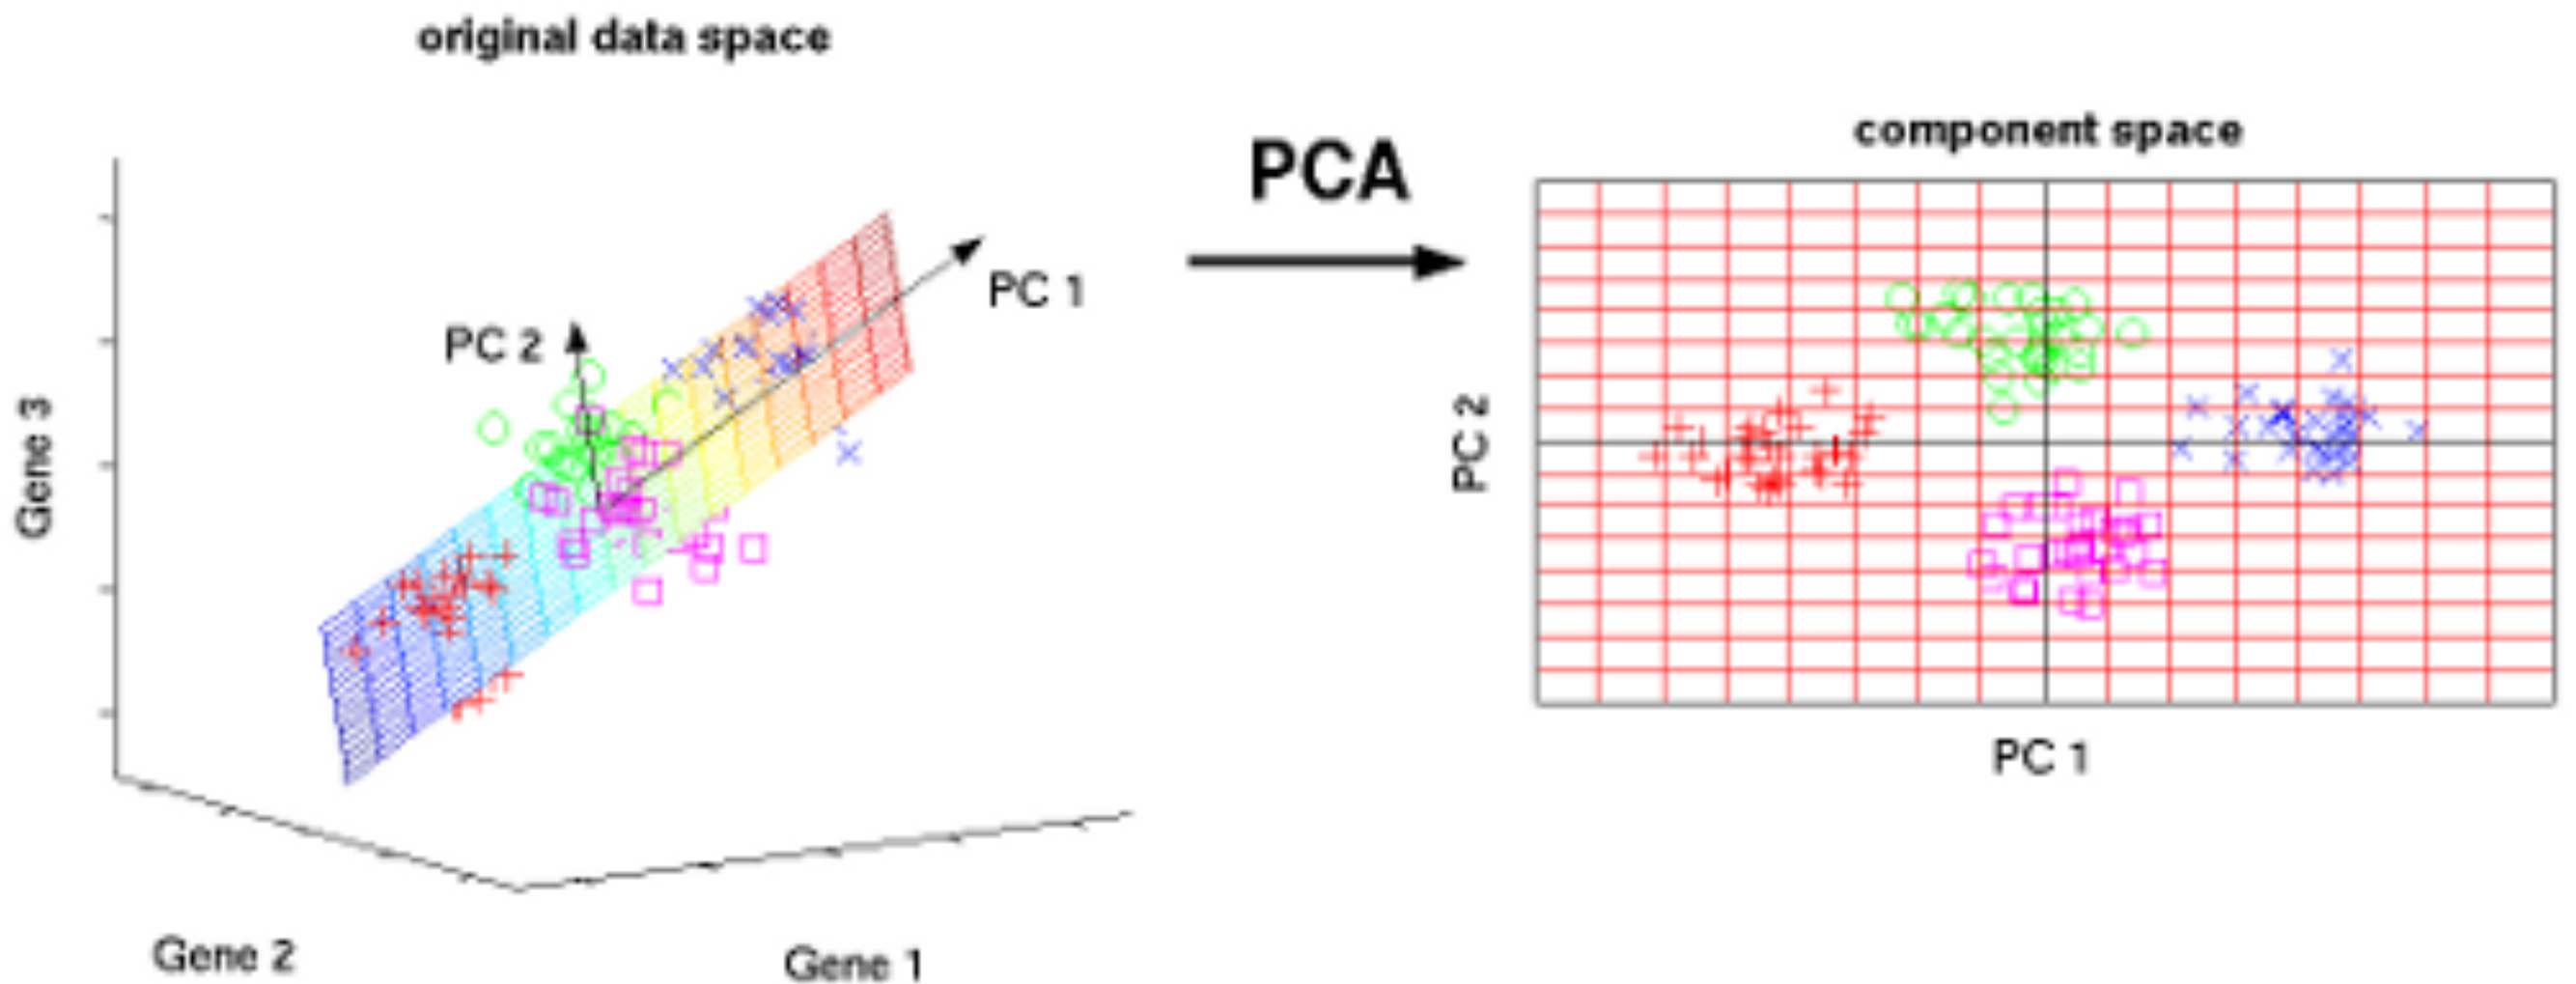
Extracting the M principal components of a vector \mathbf{x} (projection onto the M principal components)

$$\mathbf{z}(\mathbf{x}) = \mathbf{\Lambda}_M^{-\frac{1}{2}} \mathbf{U}_M^T (\mathbf{x} - \bar{\mathbf{x}})$$

- Using PCA for pre-processing gives us "normalized" vectors \mathbf{z} , i.e. uncorrelated (independent entries) and with unit variance
- We can project onto the principal directions for visualization in 2D or 3D.

Reconstruction (from the M principal components): $\mathbf{x} \approx \mathbf{U}_M (\mathbf{\Lambda}_M^{\frac{1}{2}} \mathbf{z}) + \bar{\mathbf{x}}$


Ex PCA: $D=3 \rightarrow M=2$



Source: http://www.nlpca.org/pca_principal_component_analysis.html


PCA and normalization

Original
data




Principal
components

(*intervalle* $\pm \lambda_i^{1/2}$)



Pre-processed
data


(*cov* = *I*)




We could have kept only the first
component

PCA for visualization

High-dimensional
input data




2 principal components





Ex: eigenfaces

Input data:
face images



etc ... Each example is a
d-dimensional vector

The **basis** with the 31 first eigenvectors:



<http://www.shervinemami.info/faceRecognition.html>

Kernel PCA

- The transformation learned by PCA is linear. We can get a **non-linear extension** of PCA by embedding/projecting the data into a feature space using a non-linear mapping ϕ before applying PCA

- The transformed data has the covariance matrix

$$\mathbf{C} = \frac{1}{N} \sum_{n=1}^N \phi(\mathbf{x}_n) \phi(\mathbf{x}_n)^T$$


- And the eigen-vectors in the feature space, $\mathbf{C}\mathbf{v}_i = \lambda_i \mathbf{v}_i$, can be expressed as

$$\mathbf{v}_i = \sum_{n=1}^N a_{in} \phi(\mathbf{x}_n)$$

- With a kernel k , we can use the **kernel trick** to compute dot products in the feature space **without ever having to explicitly compute the mapping**:

$$\phi(\mathbf{x}_n)^T \phi(\mathbf{x}_m) = k(\mathbf{x}_n, \mathbf{x}_m)$$

Kernel PCA: Illustration



Kernel PCA

- The kernel trick allow us to find the vector of coefficients $\mathbf{a}_i = (a_{1i}, \dots, a_{ni})^T$ representing the eigen-vectors by solving the following problem:

$$\mathbf{K}\mathbf{a}_i = \lambda_i N \mathbf{a}_i$$

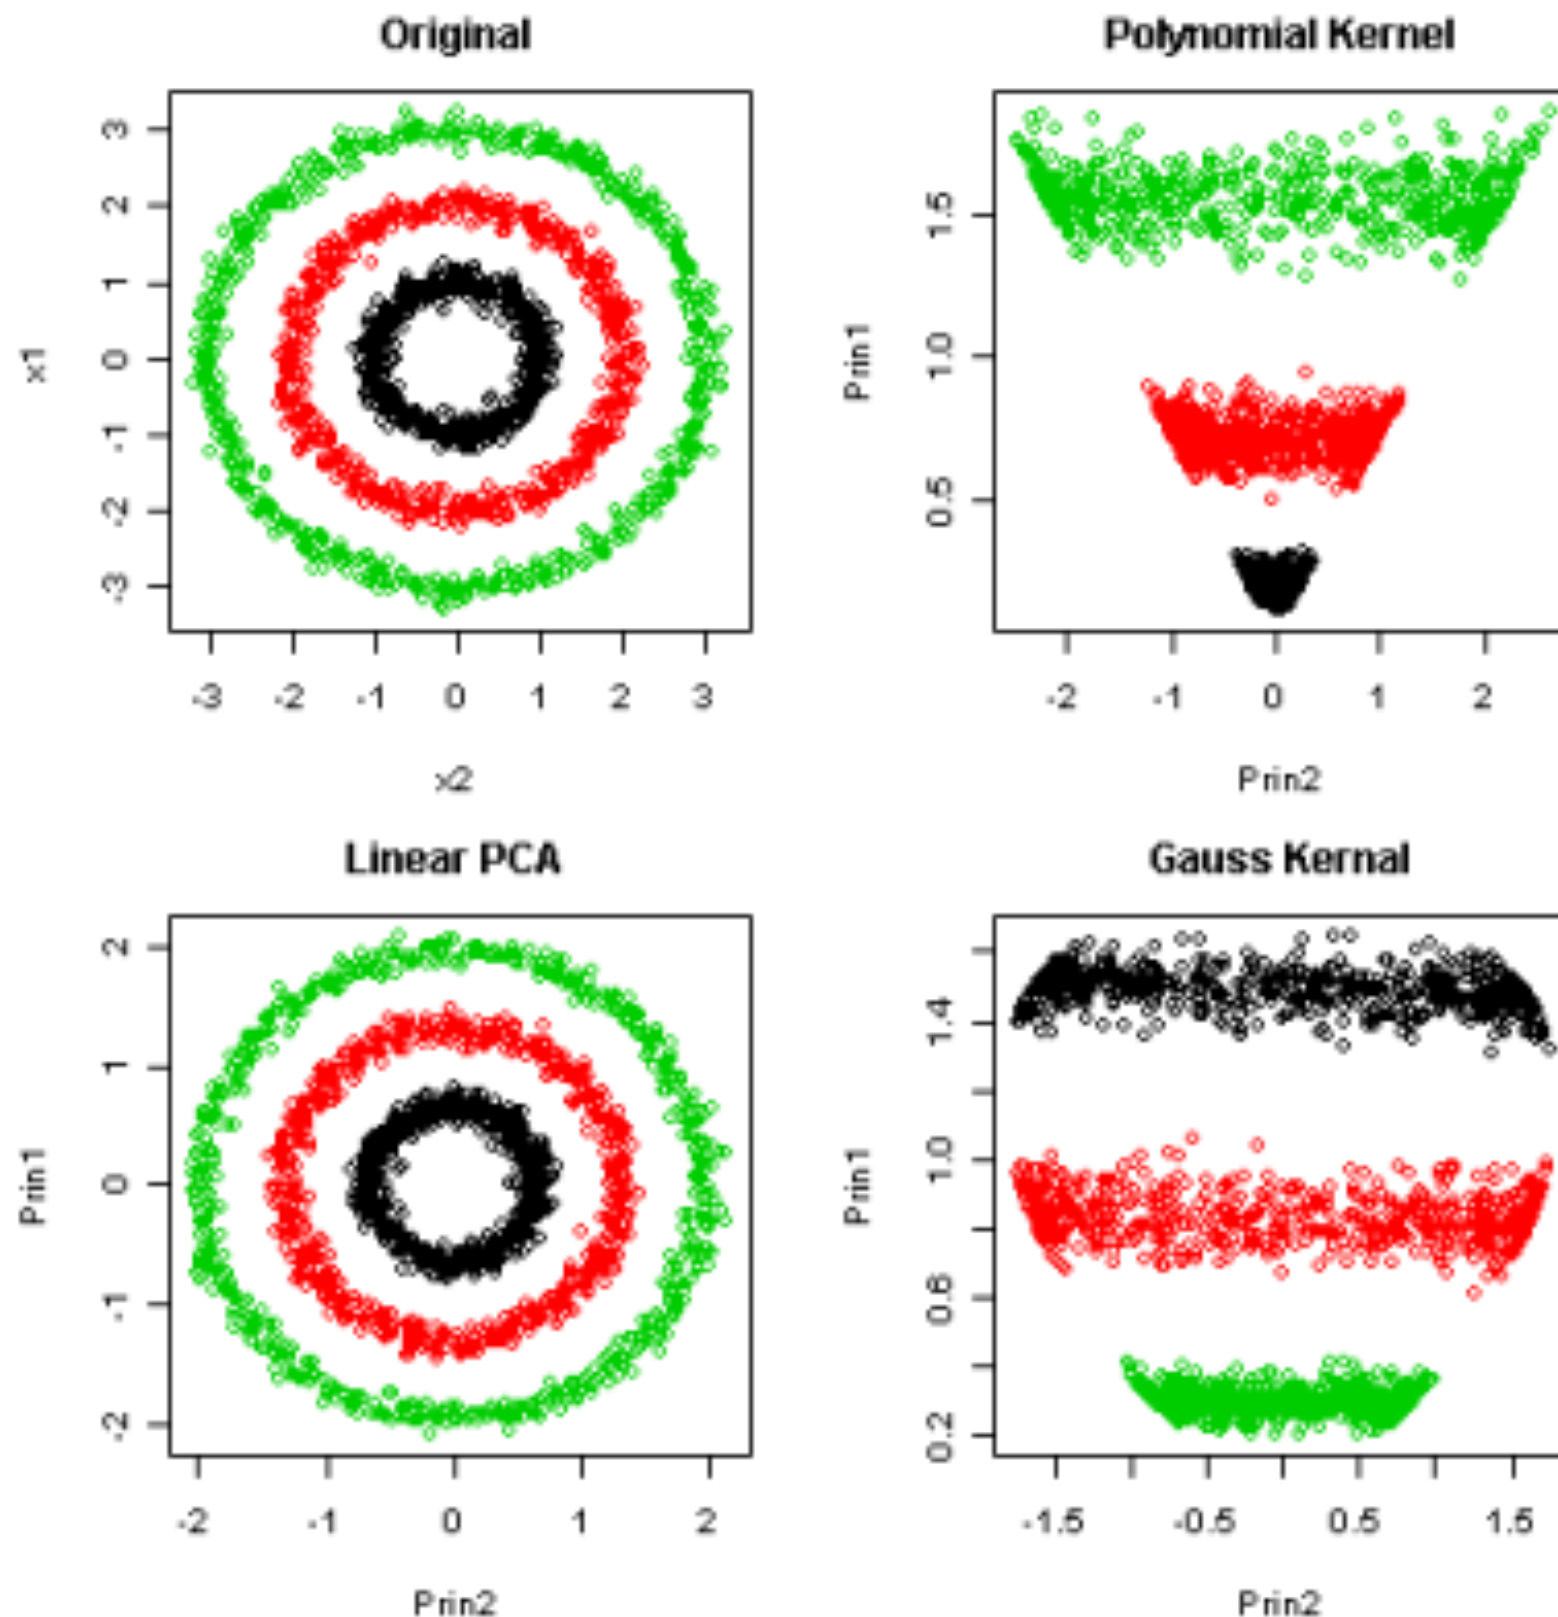
where \mathbf{K} is the Gram matrix: $\mathbf{K}_{nm} = k(\mathbf{x}_n, \mathbf{x}_m) = \phi(\mathbf{x}_n)^T \phi(\mathbf{x}_m)$

- It is then easy to compute the projections for the principal components:

$$z_i(\mathbf{x}) = \phi(\mathbf{x})^T \mathbf{v}_i = \sum_{n=1}^N a_{in} \phi(\mathbf{x})^T \phi(\mathbf{x}_n) = \sum_{n=1}^N a_{in} k(\mathbf{x}, \mathbf{x}_n)$$

- (remark: if we want to center the data in the feature space we need to use a corrected Gram matrix, c.f. Bishop 12.3)

Kernel PCA - Example




Kernel PCA: drawbacks

- Need to perform an eigen-decomposition of an $N \times N$ Gram matrix (rather than $D \times D$.)
=> This can be nice for high-dimensional data with few examples, but often $N \gg D$... Not often used with big data sets because too expensive.

AutoEncoders

- An neural network used to reproduce its input (target = input)
- The hidden layer is chosen of dimension $M < D$. This leads to **reconstruction errors**, which the training seeks to minimize.



- A **reduced-size representation** is obtained at the hidden layer.
- If the network is linear, or if it has only a hidden layer (even non-linear) it is equivalent to the PCA (the weights of the M hidden neurons define the same subspace as the PCA with M components).
- If there are several hidden layers with nonlinearities it is a **nonlinear dimensionality reduction method**.