

Density Estimation

September 28, 2018

```
In [5]: import numpy as np
import math
%pylab inline
import matplotlib.pyplot as plt

class diagonal_gaussian_parametric:
    """
    1. Implement a diagonal Gaussian parametric density estimator. It
    will have to work for data of arbitrary dimension  $d$ . As seen in
    the labs, it should have a train() method to learn the parameters
    and a method predict() which calculates the log density.
    """
    def __init__(self):
        pass

    def train(self, train_data):

        # if only one test_data is passed, add dummy dimension
        if len(np.shape(train_data)) == 1:
            train_data = np.expand_dims(train_data, axis=1)

        self.train_data = train_data

        try:
            self.n, self.d = np.shape(self.train_data)
        except:
            self.n = np.shape(self.train_data)
            self.d = 1

        self.mu = np.sum(self.train_data, axis=0) / self.n
        self.sigma = np.cov(self.train_data.T) * np.eye(self.d)

    def predict(self, test_data, nll=True):

        # if only one test_data is passed, add dummy dimension
        if len(np.shape(test_data)) == 1:
            test_data = np.expand_dims(test_data, axis=1).T
```

```

self.test_data = test_data
n_inputs = np.shape(test_data)[0]
densities = np.zeros(n_inputs)

sigma_inv = np.linalg.inv(self.sigma)
sigma_det = np.linalg.det(self.sigma)

normalizer = 1 / ((2* np.pi)**(self.d / 2) * np.sqrt(sigma_det))

# we treat each input test_data independently
for i in range(n_inputs):

    diff = self.test_data[i, :] - self.mu
    exponent = (-0.5) * (diff).T.dot(sigma_inv).dot(diff)

    if nll:
        p = -np.log(normalizer)*exponent
    else:
        p = normalizer*np.exp(exponent)

    # handle edge case where p(x)=0
    if p == 0:
        p = np.finfo(float).eps

    densities[i] = p

return(densities)

class parzen_density_estimator:
    """
    2. Implement a Parzen density estimator with an isotropic Gaussian
    kernel. It will have to work for data of arbitrary dimension d.
    Likewise it should have a train() method and a predict() method
    that computes the log density.
    """
    def __init__(self):
        pass

    def train(self, train_data, sigma=0):
        self.train_data = np.expand_dims(train_data, axis=1)

        try:
            self.n, self.d = np.shape(self.train_data)
        except:
            self.n = np.shape(self.train_data)
            self.d = 1

```

```

if sigma == 0:
    self.sigma = np.std(self.train_data) # std because isotropic Gaussian
else:
    self.sigma = sigma

def predict(self, test_data, nll=True):

    # if only one test_data is passed, add dummy dimension
    if len(np.shape(test_data)) == 1:
        test_data = np.expand_dims(test_data, axis=1).T

    self.test_data = test_data
    n_inputs = np.shape(self.test_data)[0]
    n_train = np.shape(self.train_data)[0]
    densities = np.zeros(n_inputs)
    normalizer = 1/((2*np.pi*self.sigma)**(self.d / 2))

    for i in range(n_inputs):

        # calculate avg dist b/t this training point and all test points
        p = 0
        for j in range(n_train):

            # for gaussian kernel we use euclidean distance
            distance = np.sum((self.test_data[i, :] - self.train_data[j, :])**2)
            exponent = (-0.5) * (distance / self.sigma)

            if nll:
                p += -np.log(normalizer)*exponent
            else:
                p += normalizer * np.exp(exponent)

        # handle edge case where p(x)=0
        if p == 0:
            p = np.finfo(float).eps

        # save the average kernel values across all training points
        densities[i] = p/n_train

    return(densities)

```

Populating the interactive namespace from numpy and matplotlib

```

/opt/python/sci_36/lib/python3.6/site-packages/IPython/core/magics/pylab.py:160: UserWarning: J
`%matplotlib` prevents importing * from pylab and numpy
"\n`%matplotlib` prevents importing * from pylab and numpy"

```

0.1 1D Densities

From the Iris dataset examples, choose a subset corresponding to one of the classes (of your choice), and one of the characteristic features, so that we will be in dimension $d = 1$ and produce a single graph (using the plot function) including:

- the data points of the subset (displayed on the x axis).
- a plot of the density estimated by your parametric Gaussian estimator.
- a plot of the density estimated by the Parzen estimator with a hyper-parameter (standard deviation) too small.
- a plot of the density estimated by the Parzen estimator with the hyper-parameter being a little too big.
- a plot of the density estimated by the Parzen estimator with the hyper-parameter that you consider more appropriate. Use a different color for each plot, and provide your graph with a clear legend.
- Explain how you chose your hyper-parameter .

```
In [6]: # load data
iris = np.loadtxt("iris.txt")
iris_subset = iris[:,0]
n = len(iris_subset)

# general plotting variables
n_bins = 100
axes_min = np.min(iris_subset)-1
axes_max = np.max(iris_subset)+1
alpha=0.75

small_sig = 0.001
large_sig = 4
moyen_sig = 0.25

x = np.atleast_2d(np.linspace(axes_min, axes_max, n_bins)).T

# (a) -- raw data
plt.figure(figsize=(12, 8))
plt.scatter(iris_subset, np.zeros(len(iris_subset)))

# (b) -- parametric model
flower = diagonal_gaussian_parametric()
flower.train(iris_subset)
density = flower.predict(x, nll=False)
plt.plot(x, density, color='black', alpha=alpha)

# (c) -- nonparametric small sigma
flower = parzen_density_estimator()
flower.train(iris_subset, sigma=small_sig)
density2 = flower.predict(x, nll=False)
plt.plot(x, density2, color='blue', alpha=alpha)
```

```

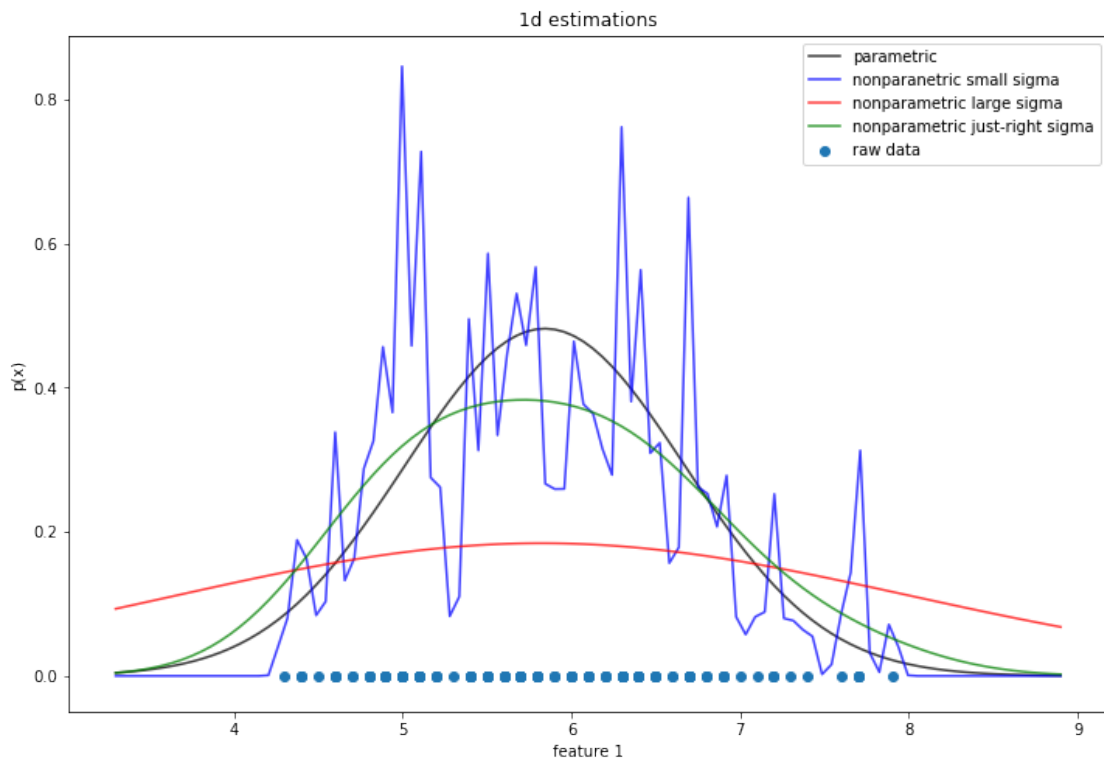
# (d) -- nonparametric large sigma
flower = parzen_density_estimator()
flower.train(iris_subset, sigma=large_sig)
density = flower.predict(x, nll=False)
plt.plot(x, density, color='red', alpha=alpha)

# (e) -- nonparametric just-right sigma
flower = parzen_density_estimator()
flower.train(iris_subset, sigma=moyen_sig)
density = flower.predict(x, nll=False)
plt.plot(x, density, color='green', alpha=alpha)

# legend
plt.legend(['parametric',
            'nonparanetric small sigma',
            'nonparametric large sigma',
            'nonparametric just-right sigma',
            'raw data'])

plt.title('1d estimations')
plt.xlabel('feature 1')
plt.ylabel('p(x)')
plt.show()

```



We chose sigma based on how it looked on the plot. The very large sigma (4) was chosen to encompass lots of empty feature space and the very small sigma (0.001) was chosen to be very sensitive to each data point. For the 'just right' sigma (0.25), we chose a value that mostly covered the data distribution without it encompassing too much empty feature space.

0.2 2D Densities

Now add a second characteristic feature of Iris, in order to have entries in $d = 2$ and produce 4 plots, each displaying the points of the subset of the data (with the plot function), and the contour lines of the density estimated (using the contour function):

- by the diagonal Gaussian parametric estimator.
- by the Parzen estimator with the hyper-parameter (standard deviation) being too small.
- by the Parzen estimator with the hyper-parameter being a little too big.
- by the Parzen estimator with the hyper-parameter that you consider more appropriate.
- Explain how you chose your hyper-parameter

```
In [7]: def make_dimensions(x_min, x_max, y_min, y_max, n_bins):
        X = np.atleast_2d(np.linspace(x_min, x_max, n_bins)).T
        Y = np.atleast_2d(np.linspace(y_min, y_max, n_bins)).T
        l = len(X) * len(Y)
        Z = np.zeros([l, 2])
        i = 0
        for x in X:
            for y in Y:
                Z[i, 0] = x
                Z[i, 1] = y
                i += 1

        return X, Y, Z
```

```
In [9]: # new data subset
iris_subset = iris[:,0:2]

# general plotting variables
n_bins = 100
x_min = np.min(iris_subset[:,0])-1
x_max = np.max(iris_subset[:,0])+1
y_min = np.min(iris_subset[:,1])-1
y_max = np.max(iris_subset[:,1])+1
alpha=0.75

# plot time
f, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2, figsize=(12, 8), sharex=True,
                                             sharey=True)

# add raw data
```

```

ax1.scatter(iris_subset[:, 0], iris_subset[:, 1])
ax2.scatter(iris_subset[:, 0], iris_subset[:, 1])
ax3.scatter(iris_subset[:, 0], iris_subset[:, 1])
ax4.scatter(iris_subset[:, 0], iris_subset[:, 1])

# (b) -- parametric model
flower = diagonal_gaussian_parametric()
flower.train(iris_subset)
X, Y, Z = make_dimensions(x_min, x_max, y_min, y_max, n_bins)
X, Y = np.meshgrid(X, Y)
density = flower.predict(Z, nll=False).reshape(n_bins, n_bins)
ax1.contour(X, Y, density.T, colors='red')
ax1.set_title('parametric model sigma={}'.format(flower.sigma[0]))
ax1.set_ylabel('feature 2')

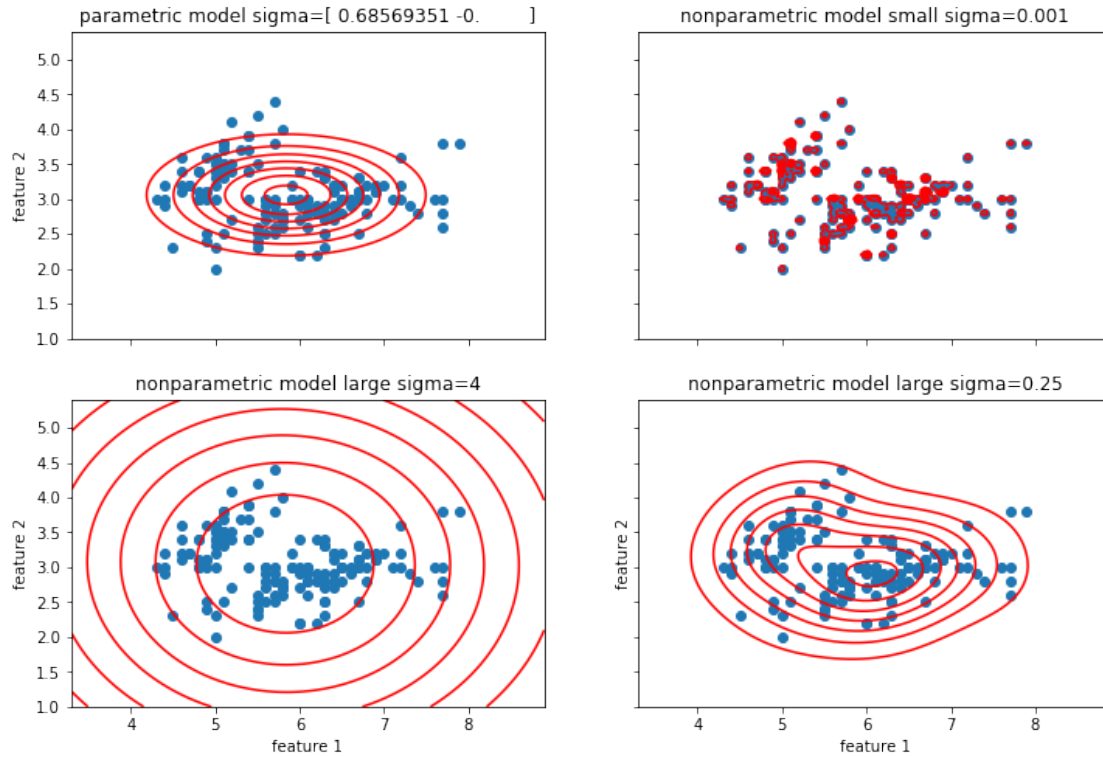
# (c) -- nonparametric model, small sigma
flower = parzen_density_estimator()
flower.train(iris_subset, sigma=small_sig)
X, Y, Z = make_dimensions(x_min, x_max, y_min, y_max, n_bins)
X, Y = np.meshgrid(X, Y)
density = flower.predict(Z, nll=False).reshape(n_bins, n_bins)
ax2.contour(X, Y, density.T, colors='red')
ax2.set_title('nonparametric model small sigma={}'.format(small_sig))

# (d) -- nonparametric model, large sigma
flower = parzen_density_estimator()
flower.train(iris_subset, sigma=large_sig)
X, Y, Z = make_dimensions(x_min, x_max, y_min, y_max, n_bins)
X, Y = np.meshgrid(X, Y)
density = flower.predict(Z, nll=False).reshape(n_bins, n_bins)
ax3.contour(X, Y, density.T, colors='red')
ax3.set_title('nonparametric model large sigma={}'.format(large_sig))
ax3.set_xlabel('feature 1')
ax3.set_ylabel('feature 2')

# (e) -- nonparametric model, just right sigma
flower = parzen_density_estimator()
flower.train(iris_subset, sigma=moyen_sig)
X, Y, Z = make_dimensions(x_min, x_max, y_min, y_max, n_bins)
X, Y = np.meshgrid(X, Y)
density = flower.predict(Z, nll=False).reshape(n_bins, n_bins)
ax4.contour(X, Y, density.T, colors='red')
ax4.set_title('nonparametric model large sigma={}'.format(moyen_sig))
ax4.set_xlabel('feature 1')
ax4.set_ylabel('feature 2')

plt.show()

```



We chose sigma based on how it looked on the plot. The very large sigma (4) was chosen to encompass lots of empty feature space and the very small sigma (0.001) was chosen to tightly wrap each data point. For the 'just right' sigma (0.25), we chose a value that mostly covered the data distribution without it encompassing too much empty feature space.